

## UNIX and C Programming (COMP1000)

# Lecture 3: Pointers

---

Updated: 13<sup>th</sup> August, 2019

Department of Computing  
Curtin University

Copyright © 2019, Curtin University  
CRICOS Provide Code: 00301J

## Textbook Reading (Hanly and Koffman)

For more information, see the weekly reading list on Blackboard.

- ▶ **Chapter 6: Pointers and Modular Programming**

Section 6.1 *also* introduces “Pointers to Files”, which you can safely ignore until lecture 5.

- ▶ **Appendix A: More About Pointers**

This is important material, even though it’s in an appendix.

For Test 1, revise everything up to, and including, next week’s lecture.

# Outline

Memory

Pointer Introduction

Using Pointers


Heap Allocation

Valgrind


# Memory

- ▶ Essentially a single giant sequence of bytes.
- ▶ Each byte has a unique, sequential “address”.
- ▶ Carved up in complex ways, into many parts.

# Memory Structure



## Code Segment




- ▶ Read-only.
- ▶ Consists of a sequence of functions.
- ▶ Each function is a sequence of machine code instructions.
- ▶ Each instruction occupies one or more bytes.
- ▶ The last instruction is the “return” instruction.

## Global Storage

- ▶ Contains statically-allocated data, including:
  - ▶ Global variables.
  - ▶ Local static variables.
  - ▶ String literals.
- ▶ Writable (not read-only), but has a fixed size

## Global Storage — Example



```
#include <stdio.h>

double globalVar = 1.5;


void printNext(void)
{
    static int count = 0;
    count++;
    printf("%d\\n", count);
}
```



# The Stack

- ▶ “Last in, first out”
- ▶ Used to store:
  - ▶ Temporary/intermediate calculation results
  - ▶ Local variables (except static variables)
  - ▶ Function parameters (sometimes)
  - ▶ Function return location (i.e. where to jump back to when a function finishes)
- ▶ Grows and shrinks over time
- ▶ Grows when a function is called (when local variables are allocated)
- ▶ Shrinks when a function returns (when local variables are destroyed)

## The Stack — Example



```
int a = 10;
int b = 12;
...
f(a, b);
...

void f(int x, int y)
{
    int z = x / y;
    printf("%d\n", z);
}
```

# The Heap


- ▶ A giant pool of dynamically-allocated memory.
- ▶ No particular fixed size or structure.
- ▶ Allocated and freed as needed by the program.
- ▶ Discussed later in this lecture.

## The Registers

- ▶ Not part of main memory; found inside the CPU itself.
- ▶ Extremely small amount of storage — only a few bytes.
- ▶ Very fast.
- ▶ Used to store (depending on the compiler and if there's enough room):
  - ▶ Function parameters.
  - ▶ Temporary/intermediate calculation results.
  - ▶ Variables declared with the `register` storage class.
- ▶ Registers also keep track of:
  - ▶ The location of the stack.
  - ▶ The currently-executing instruction.
- ▶ For most programming purposes, we can usually ignore the registers.

## Pointers

- ▶ Each byte in memory has a unique, sequential “address”.
- ▶ These addresses themselves are often stored in memory.
- ▶ This is called a “pointer”.
- ▶ Where a value occupies multiple bytes (as most do), a pointer will point to the first byte (i.e. lowest memory address).



## Pointer Data Types

- ▶ Pointers are a data type, or rather a *set of data types*.
- ▶ To declare a pointer, we must also declare what sort of data it points to.

```
int* size;  
double* speed;  
char* letter;
```

- ▶ A pointer to an int and a pointer to a float are different data types.

## Pointer Declaration Syntax

**data-type \* name;**

- ▶ Declares a variable (“**name**”), storing a memory address.
- ▶ That memory address itself stores a value of type “**data-type**” (e.g. `int`, `float`, etc.).
- ▶ Spacing is meaningless. These are all equivalent:

```
int*x;
```

```
int *x; /* Widely used. */
```

```
int* x; /* Used by me (to avoid confusion). */
```

```
int * x;
```

- ▶ Later we'll see “`*x`” used like a variable, but be careful:
  - ▶ The real variable is still `x`, not `*x`.
  - ▶ The `*` has *another* meaning: “dereference”.

## Multiple Declarations

- ▶ You can declare multiple pointers on the same line:

```
float *ptr1, *ptr2, *ptr3;
```

- ▶ Notice that they each require a \*.
- ▶ The \* is part of the data type, but attaches to one name only.
- ▶ Consider this:

```
float* ptr, number;
```

ptr is a pointer, but number is just a float.



# NULL

- ▶ A special pointer value — a “pointer to nothing”.
- ▶ Useful for initialising pointers.
- ▶ Actually a preprocessor macro, defined in all the standard header files.
- ▶ Like Java’s “null”. *Not* like SQL’s “NULL”.

## Extra Note

- ▶ NULL is *often* defined to be zero (on many platforms).
- ▶ Recall that FALSE is also zero (from the previous lecture).
- ▶ This fact is often used (or misused) inside if, while and do-while statements.

# Dereferencing

## The Uses of \*


1. Multiplication (obviously);
2. Declaring a pointer (e.g. `int* ptr`); and
3. “Dereferencing” (following) a pointer.

## The Dereference Operator

- ▶ After you’ve created a pointer, you can use it to access the memory location it points to (“dereferencing”).
- ▶ Place `*` in front of a pointer variable; e.g. “`*ptr`”.
- ▶ “`*ptr`” is what `ptr` points to. If `ptr` points to an `int`, then `*ptr` is an `int`.
- ▶ You can use `*ptr` just like an ordinary variable.
  - ▶ But remember that `ptr` is the real variable.

## Dereferencing Example

```
int* pointer;  
int number;  
  
...  
/* Assume 'pointer'  
points to a valid  
memory address! */  
  
*pointer = 5;  
number = *pointer;
```



- ▶ 5 is stored at the memory address contained in pointer.
- ▶ number is assigned the same value.

## Declaration or Dereference?

- ▶ You can declare and initialise a pointer on one line:

```
int* x = y;
```

OR

```
int *x = y; /* Remember: spacing is meaningless.*/
```

- ▶ Be careful – this does *not* dereference x!
- ▶ It *creates* x, with the datatype int\* (a pointer to an integer).
- ▶ It also initialises x with the value of y.
  - ▶ y must be another pointer of the same type (int\*).
- ▶ The following is equivalent to the above:


```
int* x;  
x = y;
```

## Obtaining Pointers — The Address-Of Operator

- ▶ How do you make a pointer point to something?
- ▶ Use the `&` (address-of) operator to get the location of a variable.
- ▶ Place a `&` in front of a variable name; e.g. `&var`.
- ▶ `&var` is a pointer to `var`.

### Example

```
int var = 15;  
int* ptr = &var;
```



Now, `*ptr` is `var`!

## Referencing and Dereferencing

- ▶ The `*` and `&` operators are opposites.
- ▶ `&` gets an address, while `*` follows an address to get a value.
- ▶ `*&var` is equivalent to just `var`.
- ▶ `&*pointer` is equivalent to `pointer`.

## Order of Operators


- ▶ Consider this: `*ptr++;`
- ▶ Dereference 1st and increment 2nd?
- ▶ Or increment 1st and dereference 2nd? (Correct!)
- ▶ To avoid ambiguity/confusion:
  - ▶ Use brackets: `(*ptr)++;`
  - ▶ Or use pre-increment: `++*ptr;`

## Pointer Example

```
int var = 10;
int* ptr;

ptr = &var; /* Make ptr point to var */
*ptr = 20;  /* Set var to 20 */

printf("%d\n", var); /* Prints "20". */
```



- ▶ The line “`ptr = &var`” stores the address of `var` in `ptr`.
- ▶ Thus, `*ptr` is `var`.
- ▶ Thus, “`*ptr = 20`” actually modifies `var`.

## Another Pointer Example

```
int b = 15;
int c = 30;
int* x;
int* y;

x = &b;    /* Make x point to b */
y = x;     /* Make y point to whatever x points to */
x = &c;     /* Make x point to c */

printf("%d\n", *x);    /* Prints "30". */
printf("%d\n", *y);    /* Prints "15". */
```

At the end:

- ▶ x points to c
- ▶ y points to b



## L-values

- ▶ Consider the humble assignment statement:

```
x = y + 5;    /* Changes x to be equal to y + 5. */
```

- ▶ The right side can be a complex expression, whereas the left side is often just a variable name. . .
- ▶ . . . but not always!
- ▶ Expressions are allowed on the left, *if* they are “L-values”.
- ▶ An L-value is a value *plus* a memory location that holds it.
  - ▶ e.g. x and y. Variables have values plus memory locations.
- ▶ Non-L-values include:
  - ▶ 5 – literal numbers have no memory location.
  - ▶ y + 5 – arithmetic has no memory location.
  - ▶ (The compiler could choose to store these in memory, but only temporarily, and it may not need to at all.)

## L-values and Pointers

- ▶ The dereference (\*) operator creates an L-value.

```
int num = 42;  
int* ptr = &num;  
...  
*ptr = 64;
```

- ▶ “\*ptr” is an expression, not a variable.
  - ▶ But it’s an L-value, so it can appear on the left.
  - ▶ \*ptr has a value (42) *and* a memory address (that of num).
- ▶ The address-of (&) operator *requires* an L-value.

```
int num = 42;  
int* ptr;  
...  
ptr = &num;
```

- ▶ You cannot write &42 or &(num + 5). They have no meaning.

## Passing by Value

- ▶ Function parameters are usually “passed by value”.
- ▶ A function’s parameters are *copies* of the values passed into it.
- ▶ The function can overwrite the copies without affecting the originals.

### Example

```
double invert(double x) {  
    x = 1.0 / x;  
    return x;  
}  
...  
double a = 0.5, b;  
b = invert(a);
```

Afterwards, `a == 0.5` and `b == 2.0`.

## Passing by Reference

- ▶ Using pointers, parameters can be “passed by reference”.
- ▶ Instead of supplying a copy of the value, you supply a pointer to the original.
- ▶ The function can then change the original variable.

### Example

```
void invert(double* x) {  
    *x = 1.0 / *x;  
}  
  
...  
double a = 0.5;  
invert(&a);
```

The value of “a” changes from 0.5 to 2.0.

## Passing by Reference (2)

- ▶ Used when the value to be passed in occupies many bytes (avoids unnecessary copying).
- ▶ Used when you want the function to export *more than one* value.
- ▶ In C, arrays and strings can only be passed by reference (but we'll get to that later).

### Example — `scanf()`


- ▶ Recall that `scanf()` requires a `&` in front of its parameters.
- ▶ These parameters are exports, not imports.
- ▶ `scanf()` needs memory addresses to store values.

## Pointers to Pointers

- ▶ A pointer can point to any data type, including another pointer.
- ▶ This is called a “pointer to a pointer (to another datatype)”.
- ▶ (a “pointer to a pointer to a pointer...”, etc. are also possible.)
- ▶ We will learn more about their uses later in the unit.

### Example

```
int num = 60;  
int* s;    /* Pointer to an integer */  
int** d;   /* Pointer to a pointer */  
  
d = &s;     /* Make d point to s */  
*d = &num;  /* Make s point to num */  
**d = 30;   /* Assign 30 to num */
```



## Void Pointers

- ▶ A pointer to “void”; a generic pointer.
- ▶ Contains a valid memory address, but the type of data stored there is unspecified.
- ▶ Cannot be dereferenced (directly):

```
void assign(int* intPtr, void* voidPtr)
{
    *intPtr = 42;    /* Perfectly fine. */
    *voidPtr = 42; /* Illegal. */
}
```

- ▶ We can't write `*voidPtr`, because we don't know the datatype.
- ▶ To use `void*`, it must be *typecast* to something else.

## Void Pointers and Typecasting

- ▶ Pointers can be typecast to different kinds of pointers:

```
void assign(void* voidPtr) {  
    int* intPtr;  
    intPtr = (int*)voidPtr; /* Typecast & assign*/  
    *intPtr = 42;  
}
```

- ▶ Here, we copy voidPtr to intPtr.
  - ▶ The typecast is highlighted.
  - ▶ Then we can access the memory.
  - ▶ We *assume* voidPtr really points to an int (exercise caution).
- ▶ We don't even need intPtr, actually:

```
void assign(void* voidPtr) {  
    *(int*)voidPtr = 42; /* Same as above. */  
}
```



## Void Pointers — Example

- ▶ Void pointers are sometimes used to handle multiple datatypes:

```
void printN(char format, void* n) {  
    if(format == 'i') {  
        /* Assume n really points to an int. */  
        printf("%d\n", *(int*)n);  
    }  
    else if(format == 'f') {  
        /* Assume n really points to a float. */  
        printf("%f\n", *(float*)n);  
    }  
}
```

- ▶ Here, printN() can be used to print ints and floats.
- ▶ The format parameter tells us how to treat the void pointer.

## Invalid Pointers


- ▶ C offers no guarantee that a pointer points to anything meaningful.
- ▶ It's your job to make sure it does.
- ▶ Newer languages (like Java) use pointers “behind the scenes”, and so protect you from misusing them.
- ▶ C offers no protection.

### Example


Say your function returns a pointer to one of its own local variables:

- ▶ The local variable will be destroyed when the function ends.
- ▶ The pointer will point to memory no longer in use.

## Invalid Pointers — Example



```
int* ptr;  
ptr = func();  
...  
  
int* func(void)  
{  
    int var = 33;  
    return &var;  
}
```



## Heap Memory

- ▶ So far we've dealt with the stack and global storage.
- ▶ The “heap” is a much larger (and more flexible) pool of memory.
- ▶ Your functions can return pointers to heap-allocated memory.
  - ▶ You don't have to allocate and use the memory in the same function.
  - ▶ Can't be done with stack-based variables.
- ▶ Mostly useful for arrays (lecture 4) and structs (lecture 6).

# Allocation

- ▶ All memory must be “allocated” before it can be used.
- ▶ Stack memory is allocated simply by declaring a local variable.
- ▶ Heap memory is allocated by calling the `malloc()` function.
- ▶ The OS finds a block of memory of a suitable size, and grants your program permission to access it.
- ▶ This memory *stays* allocated until you “free” it — it doesn’t disappear when a function ends.

## Deallocation / Freeing

- ▶ Newer languages (like Java) have “garbage collection”.
  - ▶ When a block of heap memory no longer has any pointers to it, it is automatically freed.
- ▶ C **does not** have garbage collection!
- ▶ You must explicitly deallocate heap memory when you're done with it.

## The Operating System

- ▶ When a program ends, all its allocated memory *is* forceably freed by the OS.
- ▶ For large or long-running programs, this doesn't help you.
- ▶ Don't rely on it!

## Memory Leaks

- ▶ If you forget to deallocate heap memory in C, you get a “memory leak”.
- ▶ This is often an invisible error.
- ▶ Your program may appear to work perfectly.
- ▶ However, it may consume much more memory than needed (and you may run out of memory).
- ▶ There are tools — like `valgrind` — to help detect memory leaks.

## Heap Allocation in C

- ▶ Heap allocation (and de-allocation) is done through specific functions.
- ▶ Mainly uses the `stdlib.h` library.
- ▶ Some generic memory-related functions are found in `string.h`.
  - ▶ (These include `memset()` and `memcpy()`, which we'll cover next week.)



## The sizeof Operator

- ▶ Reports the size (in bytes) of any data type.
- ▶ A C language construct (*not* a function).
- ▶ On a typical 32-bit machine:
  - `sizeof(char) == 1`
  - `sizeof(short) == 2`
  - `sizeof(int) == 4`
  - `sizeof(float) == 4`
  - `sizeof(double) == 8`
  - `sizeof(void*) == 4`
  - `sizeof(short**) == 4`
- ▶ These sizes may change across different hardware — that's why `sizeof` is useful!
- ▶ A special unsigned integer type “`size_t`” represents memory sizes:

```
size_t numBytes = sizeof(long);
```

## Heap Allocation — `malloc()`

- ▶ The `malloc()` function allocates a block of heap memory.
- ▶ Takes an `int` parameter — the size in bytes.
- ▶ Returns a `void*` pointing to the newly-allocated memory.
- ▶ Returns `NULL` if the memory could not be allocated.

### Usage

Since `malloc()` doesn't know (or care) what you want to store in the memory:

- ▶ Use `sizeof` to determine the number of bytes to allocate.
- ▶ Typecast the returned `void` pointer to the appropriate pointer type.

## malloc() Example

To *dynamically* allocate storage for various values:

```
#include <stdlib.h>

...

int* integer = (int*)malloc(sizeof(int));
float* real = (float*)malloc(sizeof(float));
double* bigReal = (double*)malloc(sizeof(double));
```

- ▶ “sizeof(int)” — the number of bytes to allocate.
- ▶ “(int\*)” — typecast to an int pointer.
- ▶ Once allocated, you can use \*integer, \*real and \*bigReal just like “ordinary” variables.

## Deallocation — `free()`

- ▶ Every block of memory allocated with `malloc()` must eventually be “freed”.
- ▶ The `free()` function takes a pointer to the block, and frees it.
- ▶ Returns nothing.
- ▶ Thereafter, the pointer is invalid.
- ▶ Don't free memory before you're finished with it!

### NULL

- ▶ It's good practice to set a freed pointer to `NULL`.
- ▶ You should do this immediately after a call to `free()`.
  - ▶ (Unless the pointer variable itself is about to disappear too.)

## free() — Example

```
#include <stdlib.h>

int main(void) {
    double* real = (double*)malloc(sizeof(double));

    ... /* Use the memory */

    free(real);
    real = NULL;

    ... /* Do something else afterwards */
}
```

- ▶ You don't need to tell `free()` how big the block is.
- ▶ Make sure you set the pointer to `NULL` afterwards.

# Valgrind

- ▶ It's easy to make mistakes with memory in C:
  - ▶ Using uninitialised values,
  - ▶ Accessing unallocated memory,
  - ▶ Failing to free memory,
  - ▶ Freeing the same memory more than once,
  - ▶ Losing track of allocated memory (memory leaks).
- ▶ The valgrind tool helps you find these.
- ▶ valgrind works with compiled programs.
- ▶ It detects memory errors while a program is running.
- ▶ To use it, type:

```
[user@pc]$ valgrind [options] ./program ...
```

Leave **[options]** blank for a summary of the errors.

## Debugging Information

- ▶ To get the most out of valgrind, you need “debugging information”.
- ▶ This is a compile option, which inserts extra information into the executable file.
- ▶ On the command-line:

```
[user@pc]$ gcc -g -c file.c
```

- ▶ In a Makefile:

```
CFLAGS = -Wall -pedantic -ansi -g
```

- ▶ Without this, valgrind can't give you any line numbers.

## Valgrind Output

This is valgrind detecting a memory leak:

### LEAK SUMMARY:

definitely lost: 10 bytes in 1 blocks

indirectly lost: 0 bytes in 0 blocks

possibly lost: 0 bytes in 0 blocks

still reachable: 0 bytes in 0 blocks

suppressed: 0 bytes in 0 blocks

Rerun with `--leak-check=full` to see details of  
leaked memory

Then, to find *where* the leak occurred:

```
[user@pc]$ valgrind --leak-check=full ./program ...
```



## Coming Up

- ▶ Next week's lecture will expand on pointers, covering arrays and strings.
- ▶ Once again, make sure you complete the tutorial exercises.